

CS5830

Image

Classification

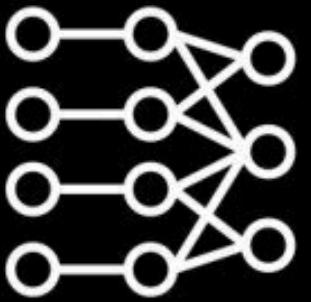
Team 09

R Prithiviraj (ME20B136)

B Srivathsan (ME20B176)

R S Harish Krishnan (ME20B081)

Overview



01



02

mlflow



03

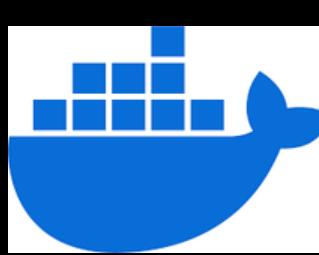
04



05



06

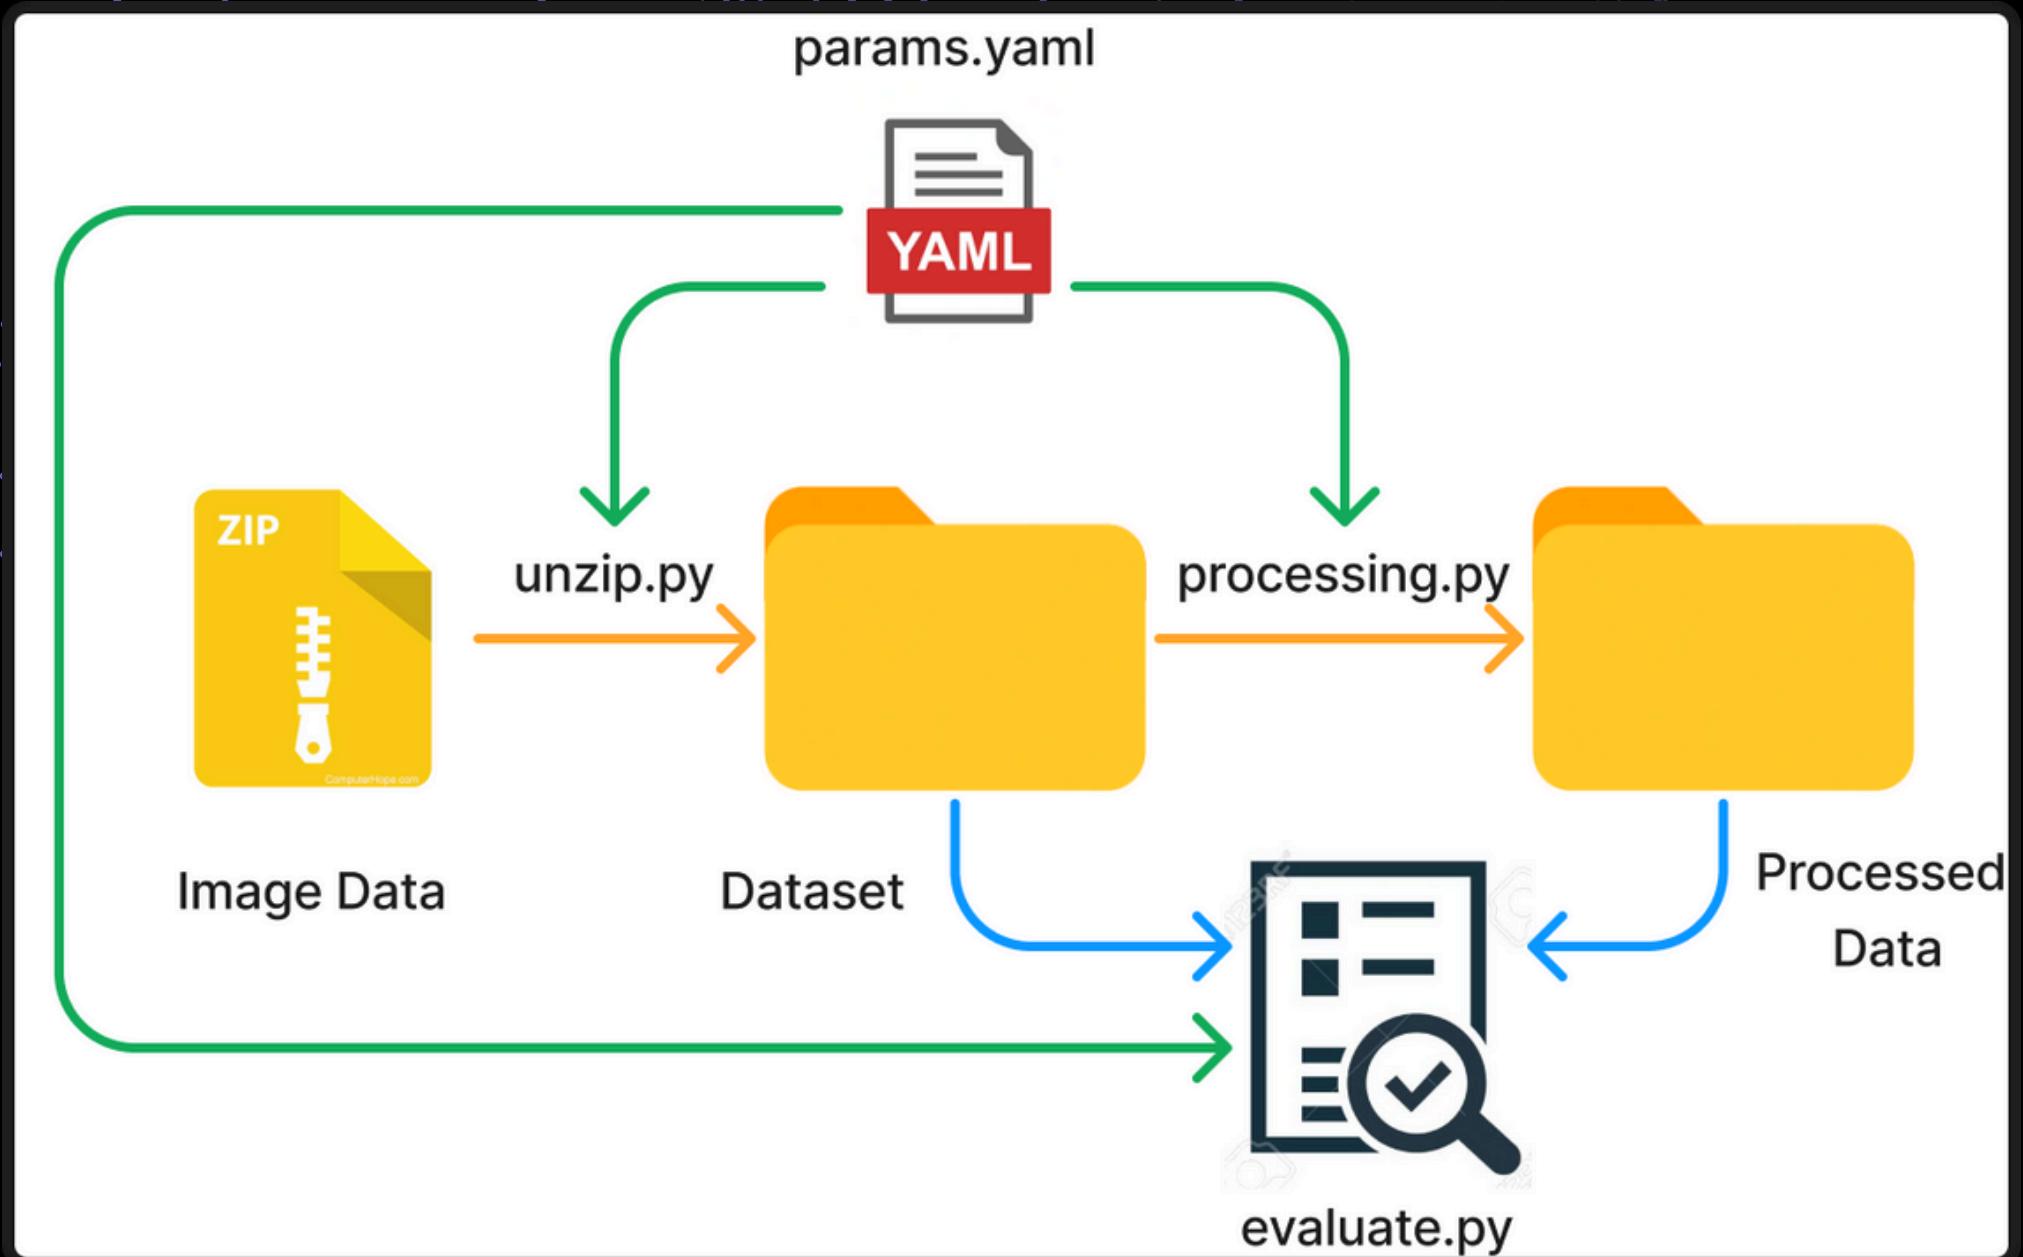


07



08

Data Processing



- 5 class dataset is available as a .zip file
- 3 stages were used for processing the data
- Stages Involved:
 - Unzipping the Dataset
 - Processing the file using **multiprocessing** as an alternative for **Spark**
 - Image Resizing
 - Mean Variance Normalization
 - Evaluate:
 - Check if the unzipped file and processed file have similar content structure
- Git was used for tracking the files and version control.
- DVC was used to pipeline the stages.

DAG

```
+-----+
| unzip |
+-----+
**          **
*          *
+-----+          *
| preprocess |          **
+-----+          **
**          **
*          *
+-----+          *
| evaluate |          **
+-----+
```

DVC Prithiviraj7R > BDL-Project Private :

+ Add a model Help

Search experiments 1 Plots Trends Compare

Experiment	Created	Message	CML	params.yaml										
				.../live/metrics.json		base		evaluate		preprocess				
				Matching	zip_file_path	ORIGINAL_DAT...	OUTPUT_PATH	PROCESSED_D...	mean	output_folder	std	test_folder	train_folder	validation_folder
<input type="checkbox"/>  main :														10 of 10 experiments (8 hidden)
<input type="checkbox"/>  BASELINE HEAD, main :	04:59 PM	 Update README.md		true	datasets.zip	datasets	eval_results	processed_data	[0.485,0.456,...]	processed_data	[0.229,0.224,...]	datasets/test	datasets/train	—
<input type="checkbox"/> bc868f3 : 	04:50 PM	 Update README.md		true	datasets.zip	datasets	eval_results	processed_data	[0.485,0.456,...]	processed_data	[0.229,0.224,...]	datasets/test	datasets/train	—
<input type="checkbox"/> a38df1c : 	04:50 PM	 Update README.md		true	datasets.zip	datasets	eval_results	processed_data	[0.485,0.456,...]	processed_data	[0.229,0.224,...]	datasets/test	datasets/train	—
<input type="checkbox"/> 83b85a2 : 	04:33 PM	 Update README.md		true	datasets.zip	datasets	eval_results	processed_data	[0.485,0.456,...]	processed_data	[0.229,0.224,...]	datasets/test	datasets/train	—
<input type="checkbox"/> 3745f32 : 	04:06 PM	 Added Prometheus and Gr...		true	datasets.zip	datasets	eval_results	processed_data	[0.485,0.456,...]	processed_data	[0.229,0.224,...]	datasets/test	datasets/train	—
<input type="checkbox"/> e495821 : 	04:00 PM	 Model Selection using MLfl...		true	datasets.zip	datasets	eval_results	processed_data	[0.485,0.456,...]	processed_data	[0.229,0.224,...]	datasets/test	datasets/train	—
<input type="checkbox"/> 96158b2 : 	03:25 PM	 Experiment run in DVC		true	datasets.zip	datasets	eval_results	processed_data	[0.485,0.456,...]	processed_data	[0.229,0.224,...]	datasets/test	datasets/train	—

Modelling

- Evaluated five CNN models with varying
 - Filter sizes
 - Number of filters
 - Batch normalization
 - Learning rate
 - Number of epochs
- Tracked performance metrics using mlflow, highlighting optimal parameter combinations for best accuracy and lowest loss.
- Results can be visualized through the mlflow UI

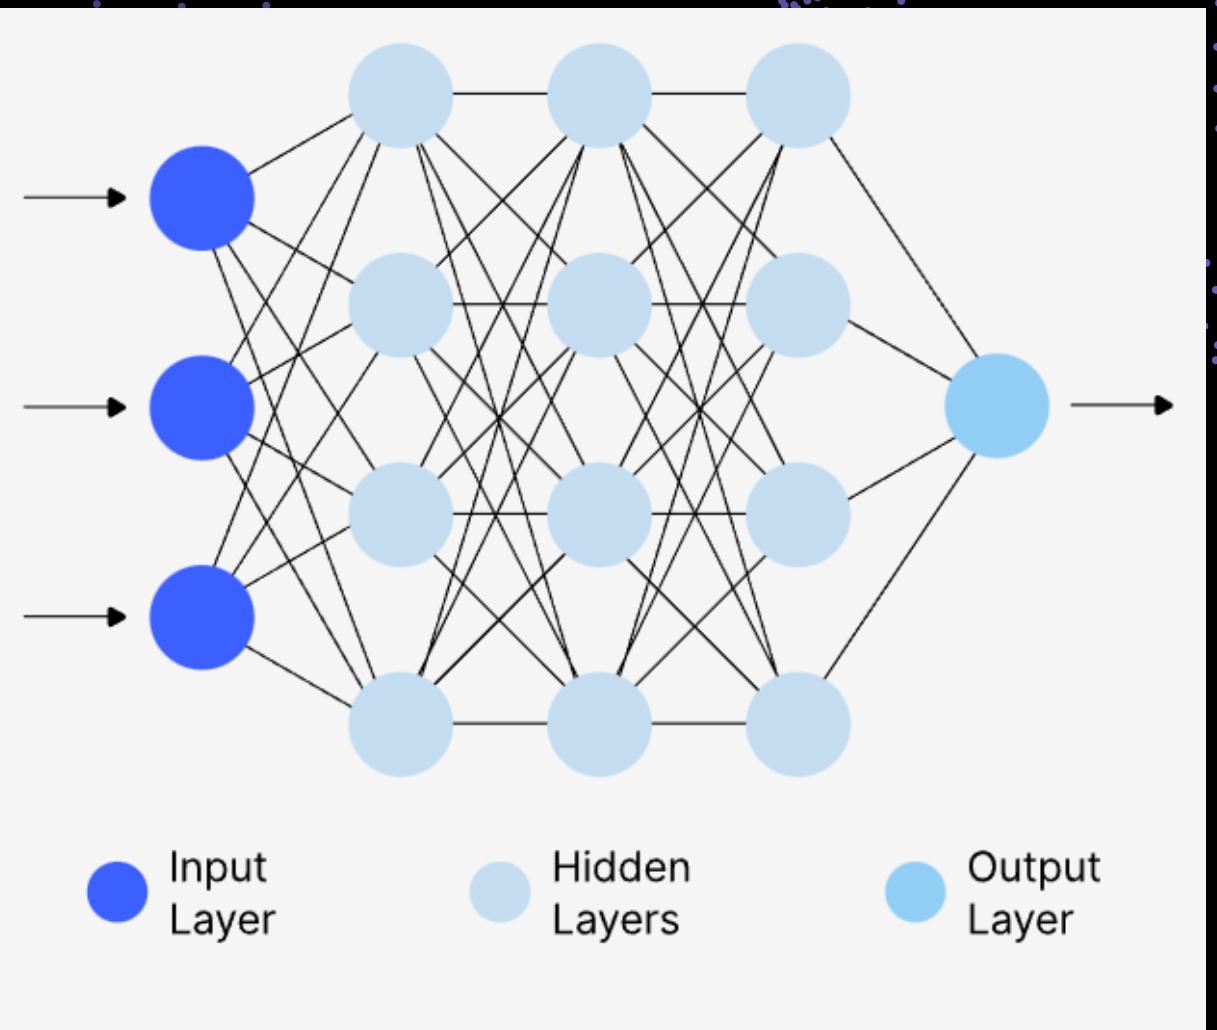
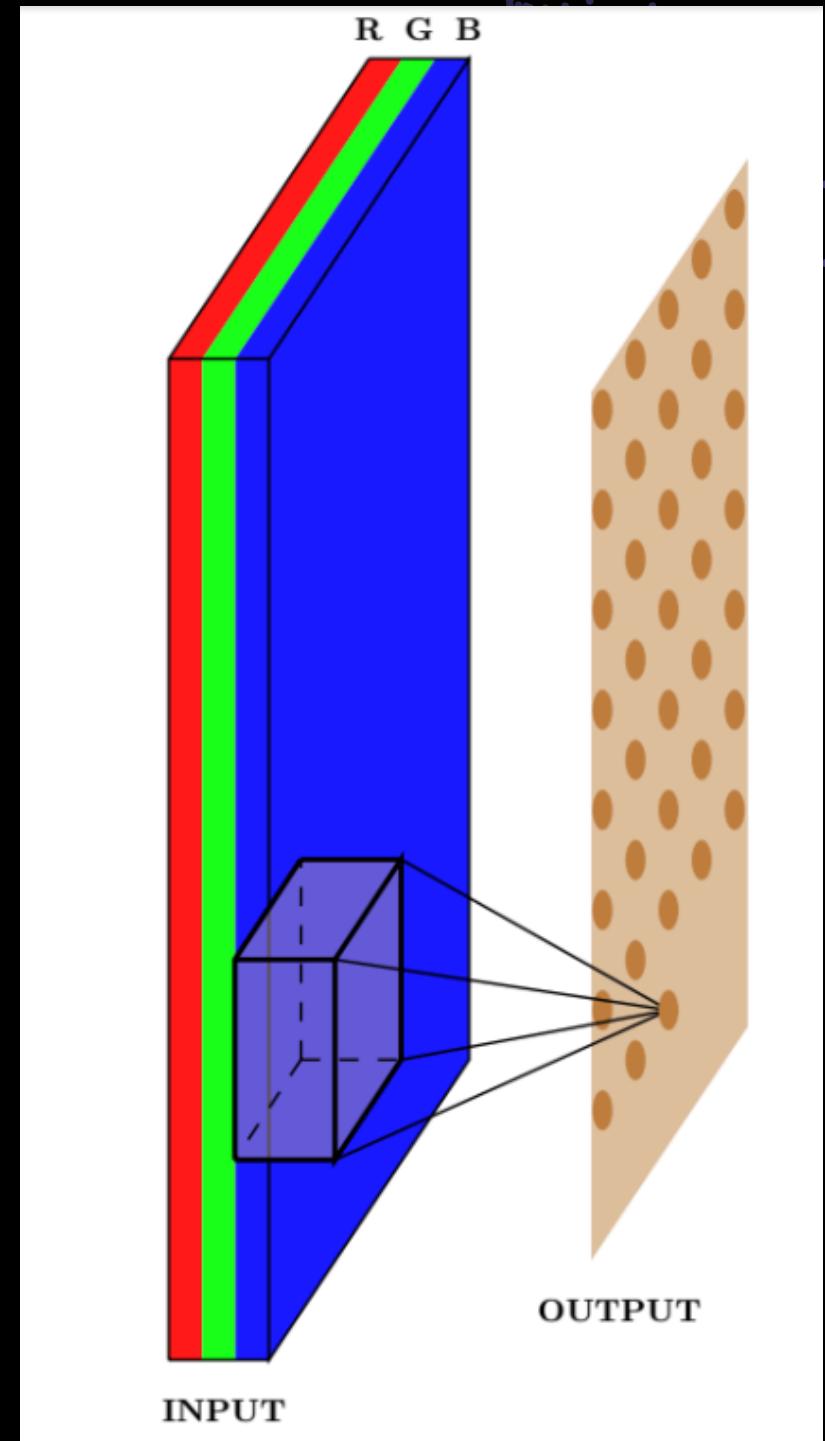


Image classification

[Provide Feedback](#)

Add Description

Share

Q metrics.rmse < 1 and params.model = "tree"



Time created

State: Active

Datasets

Sort: Created

Columns

Group by

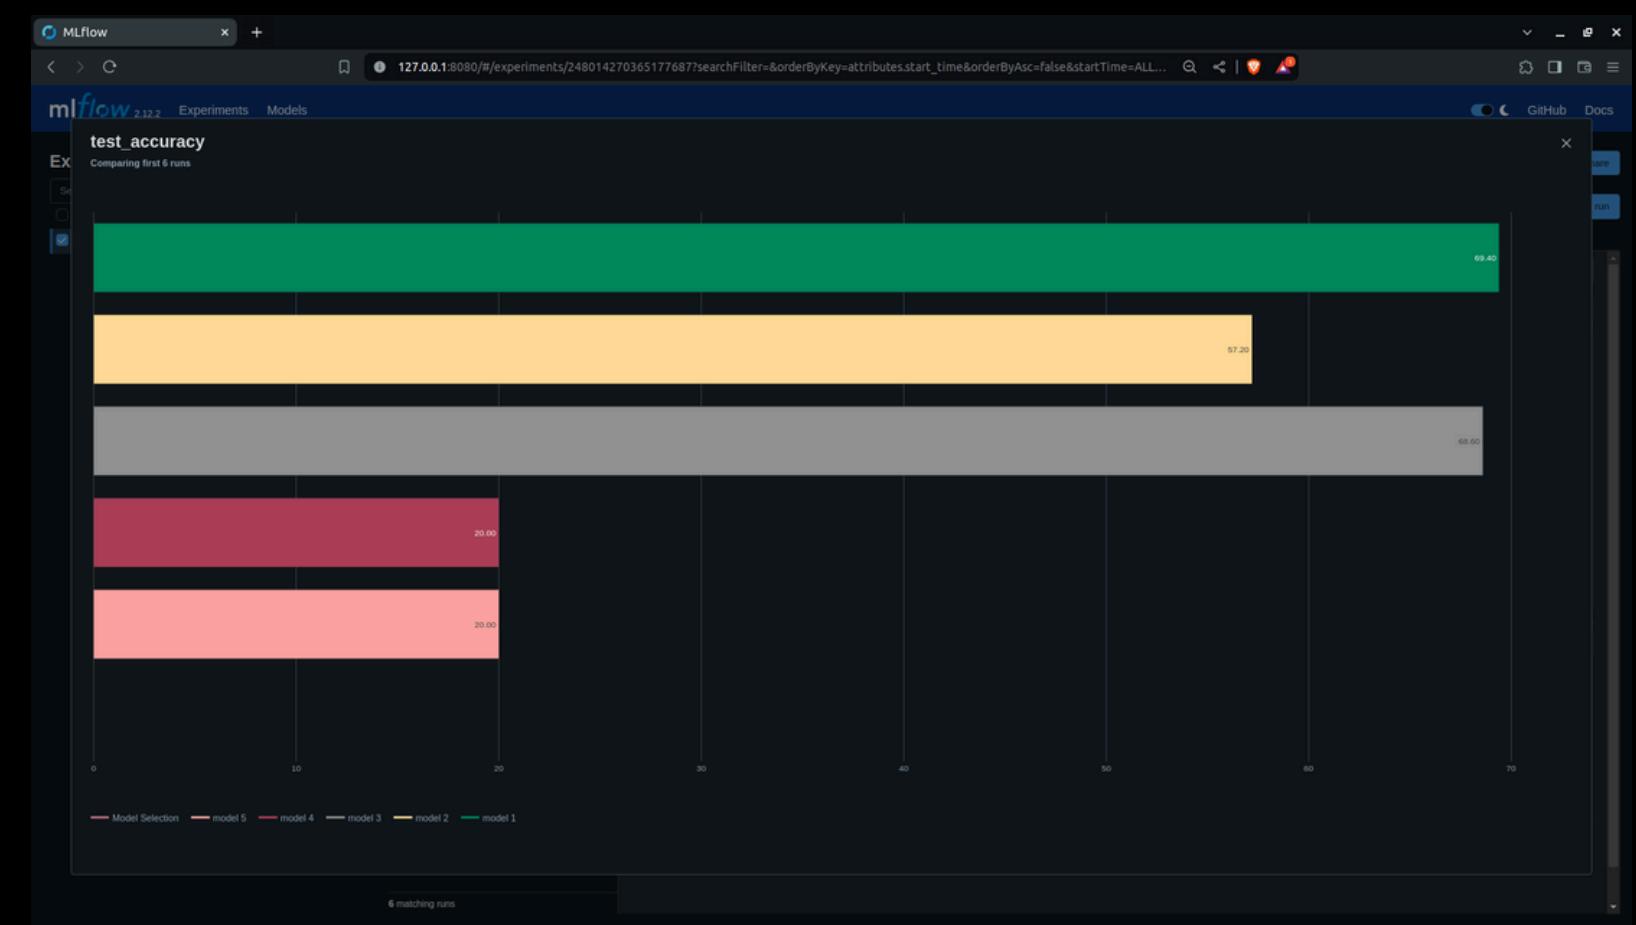
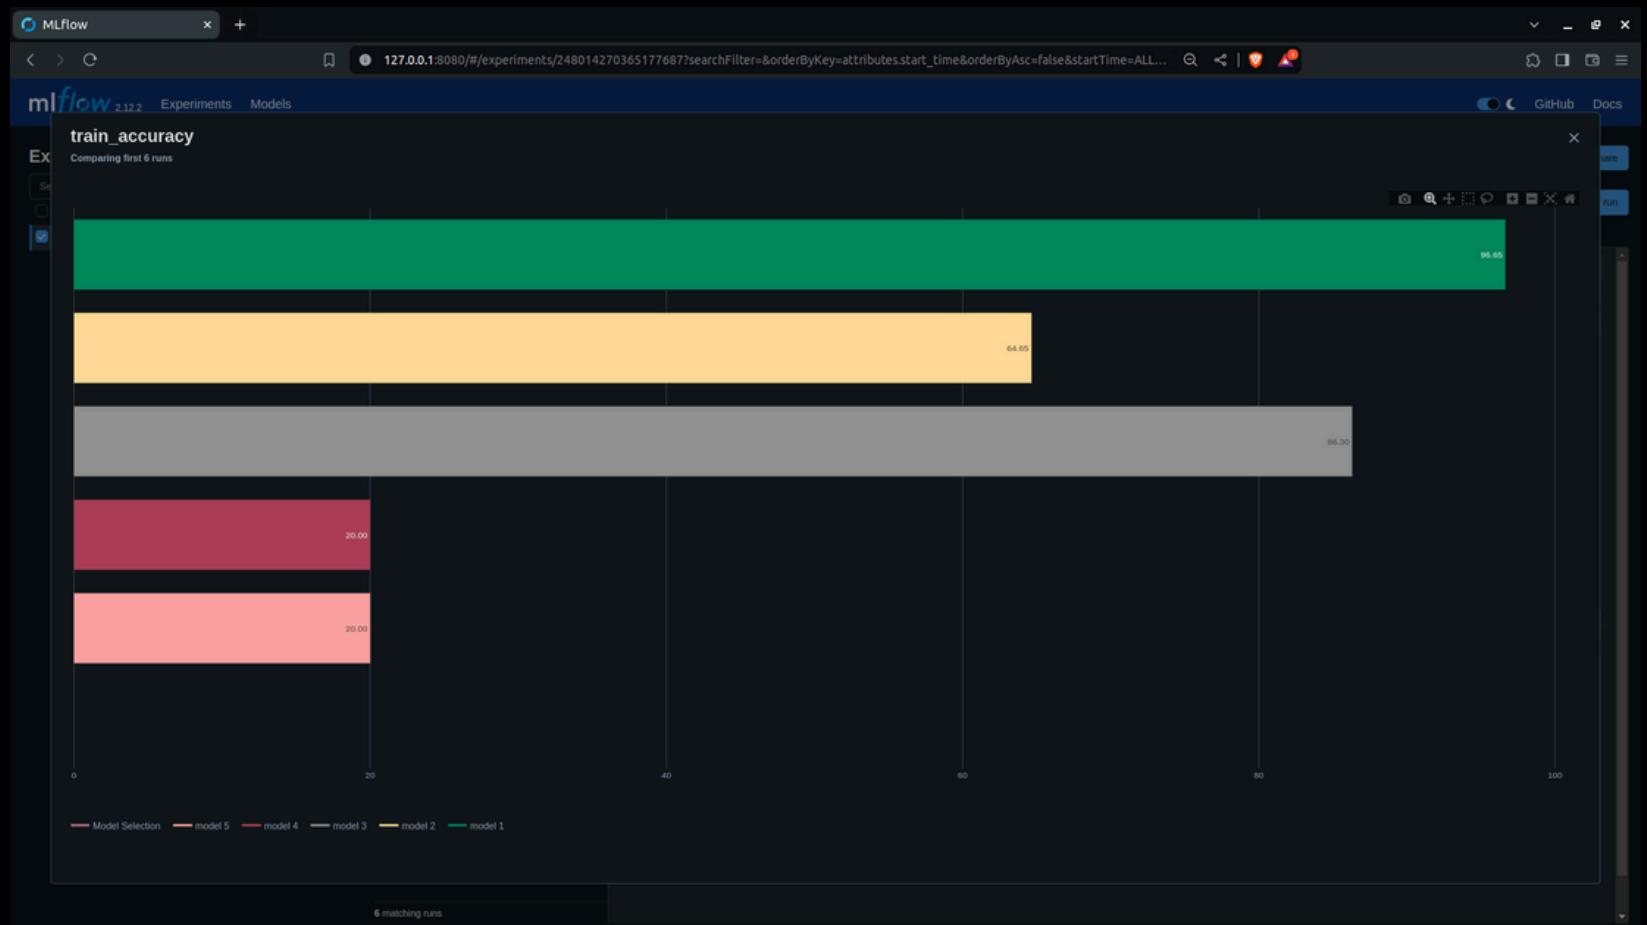


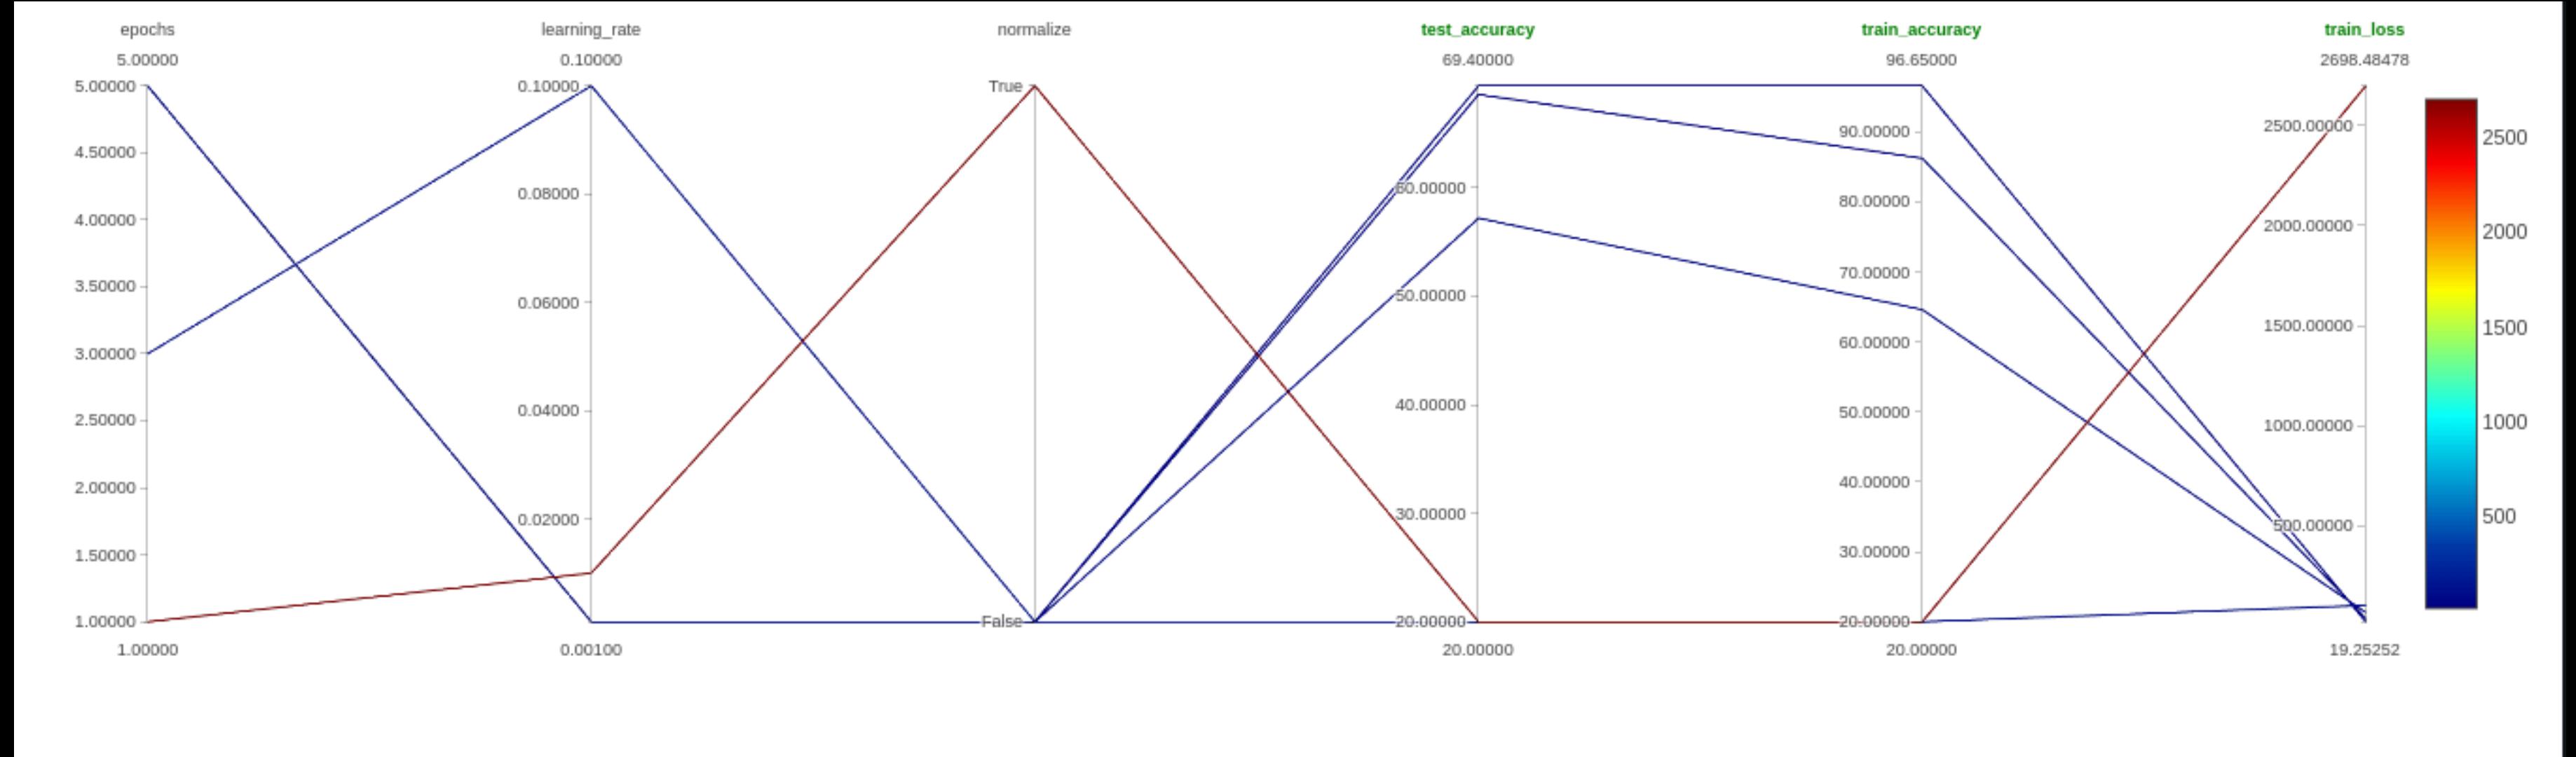
+ New run

Table Chart Evaluation Experimental

Metrics													Parameters		
	Run Name	Created	Duration	test_accuracy	train_accuracy	train_loss	conv1	conv2	epochs	learning_rate	normalize				
	Model Selection	19 minutes ago	19.2min	-	-	-	-	-	-	-	-				
	model 5	1 minute ago	1.5min	20	20	101.820723...	{'in_channe...	{'in_channe...	3	0.1	False				
	model 4	3 minutes ago	1.5min	20	20	2698.48478...	{'in_channe...	{'in_channe...	1	0.01	True				
	model 3	11 minutes ago	7.2min	68.6000000...	86.3	35.8205042...	{'in_channe...	{'in_channe...	5	0.001	False				
	model 2	17 minutes ago	4.6min	57.1999999...	64.6499999...	65.8338003...	{'in_channe...	{'in_channe...	5	0.001	False				
	model 1	19 minutes ago	2.1min	69.3999999...	96.65	19.2525206...	{'in_channe...	{'in_channe...	5	0.001	False				

+
Show more columns
(7 total)

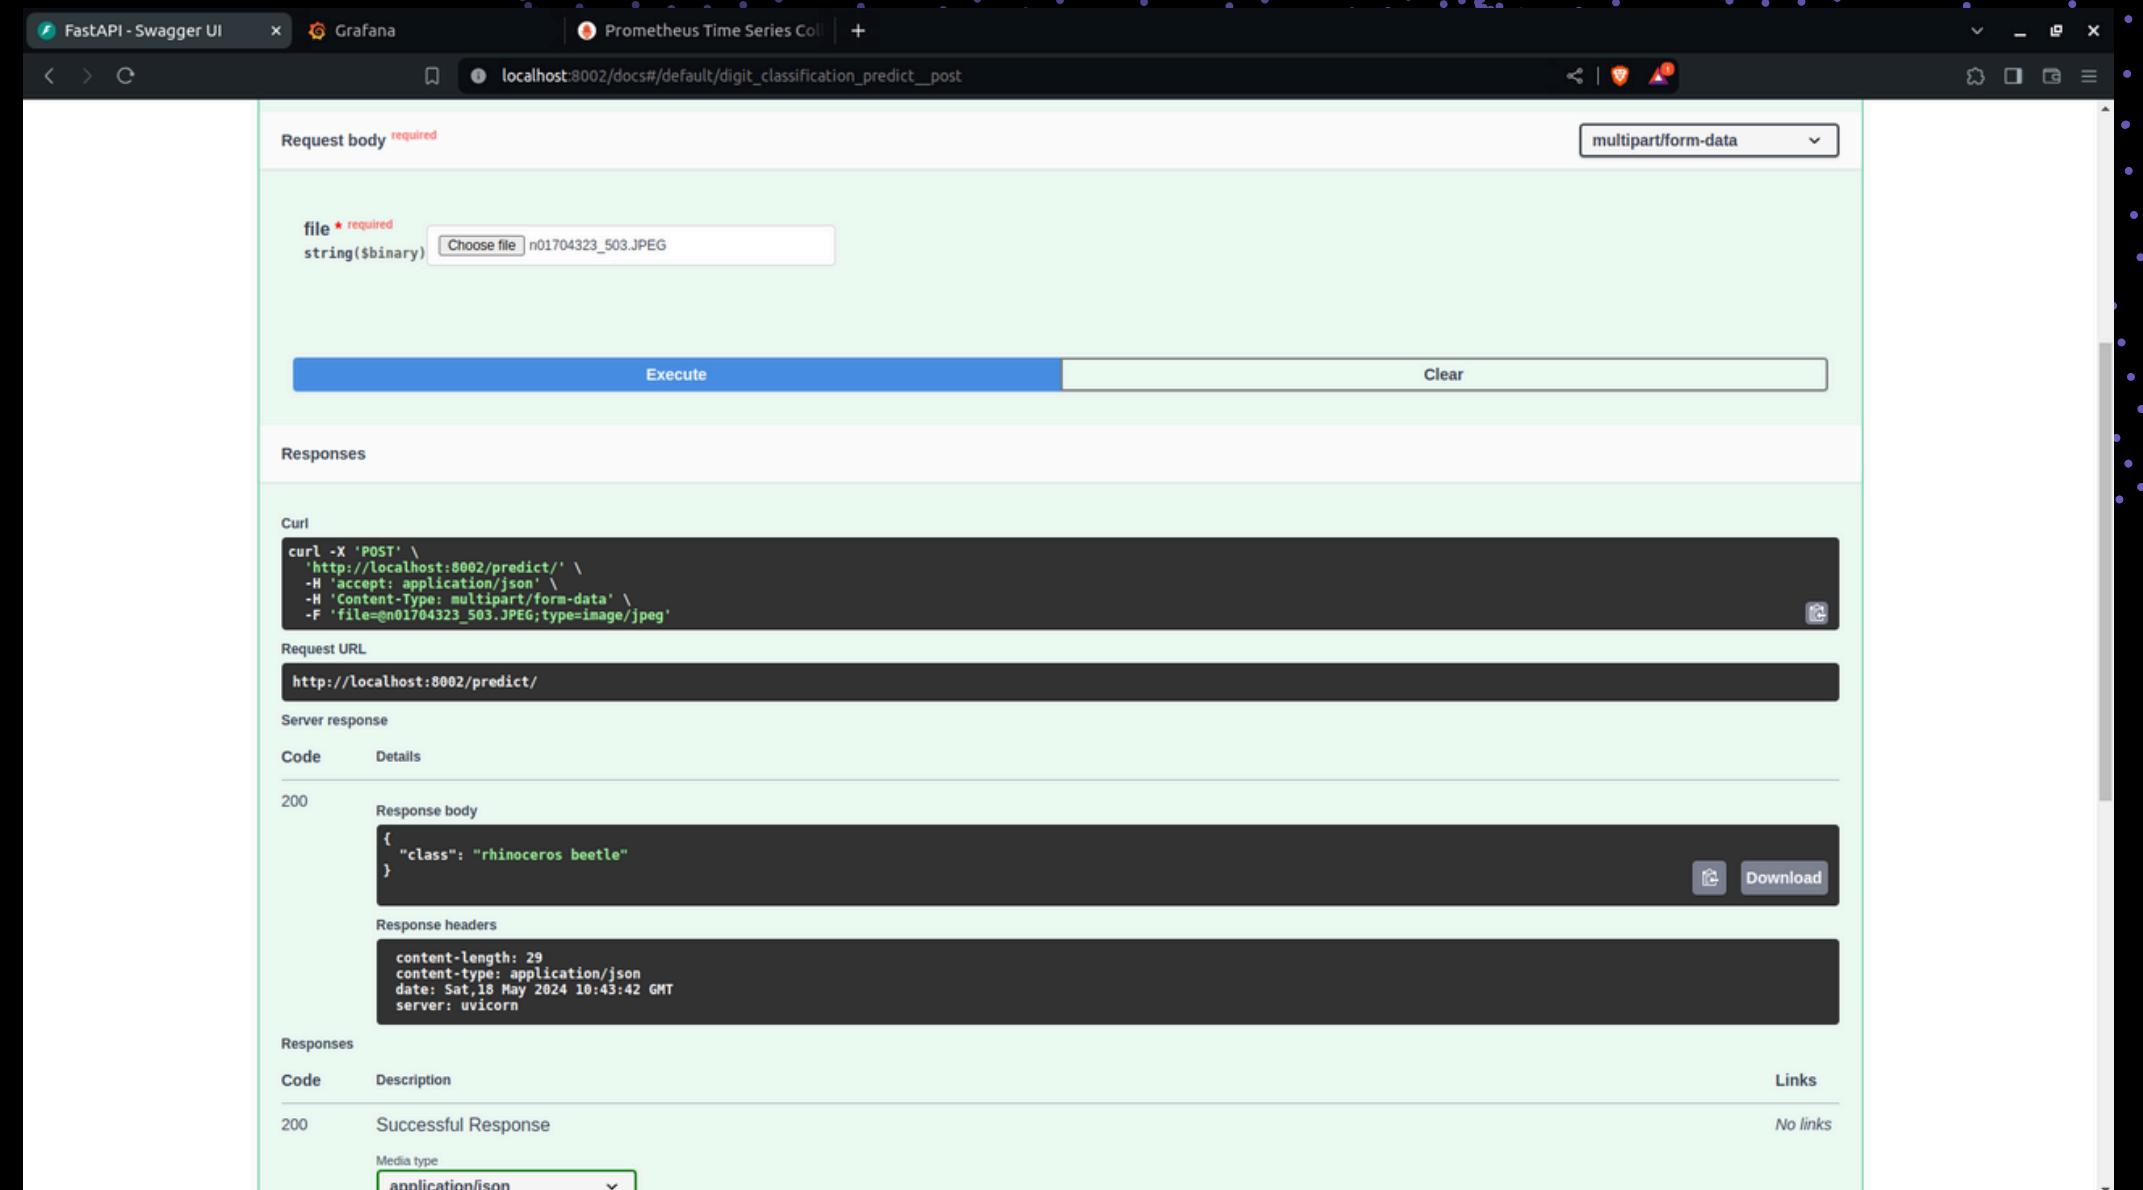


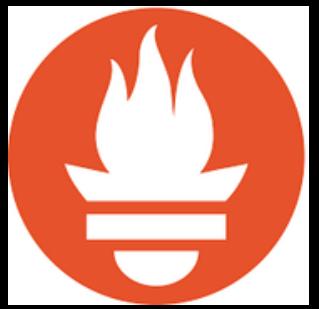




Model Deployment

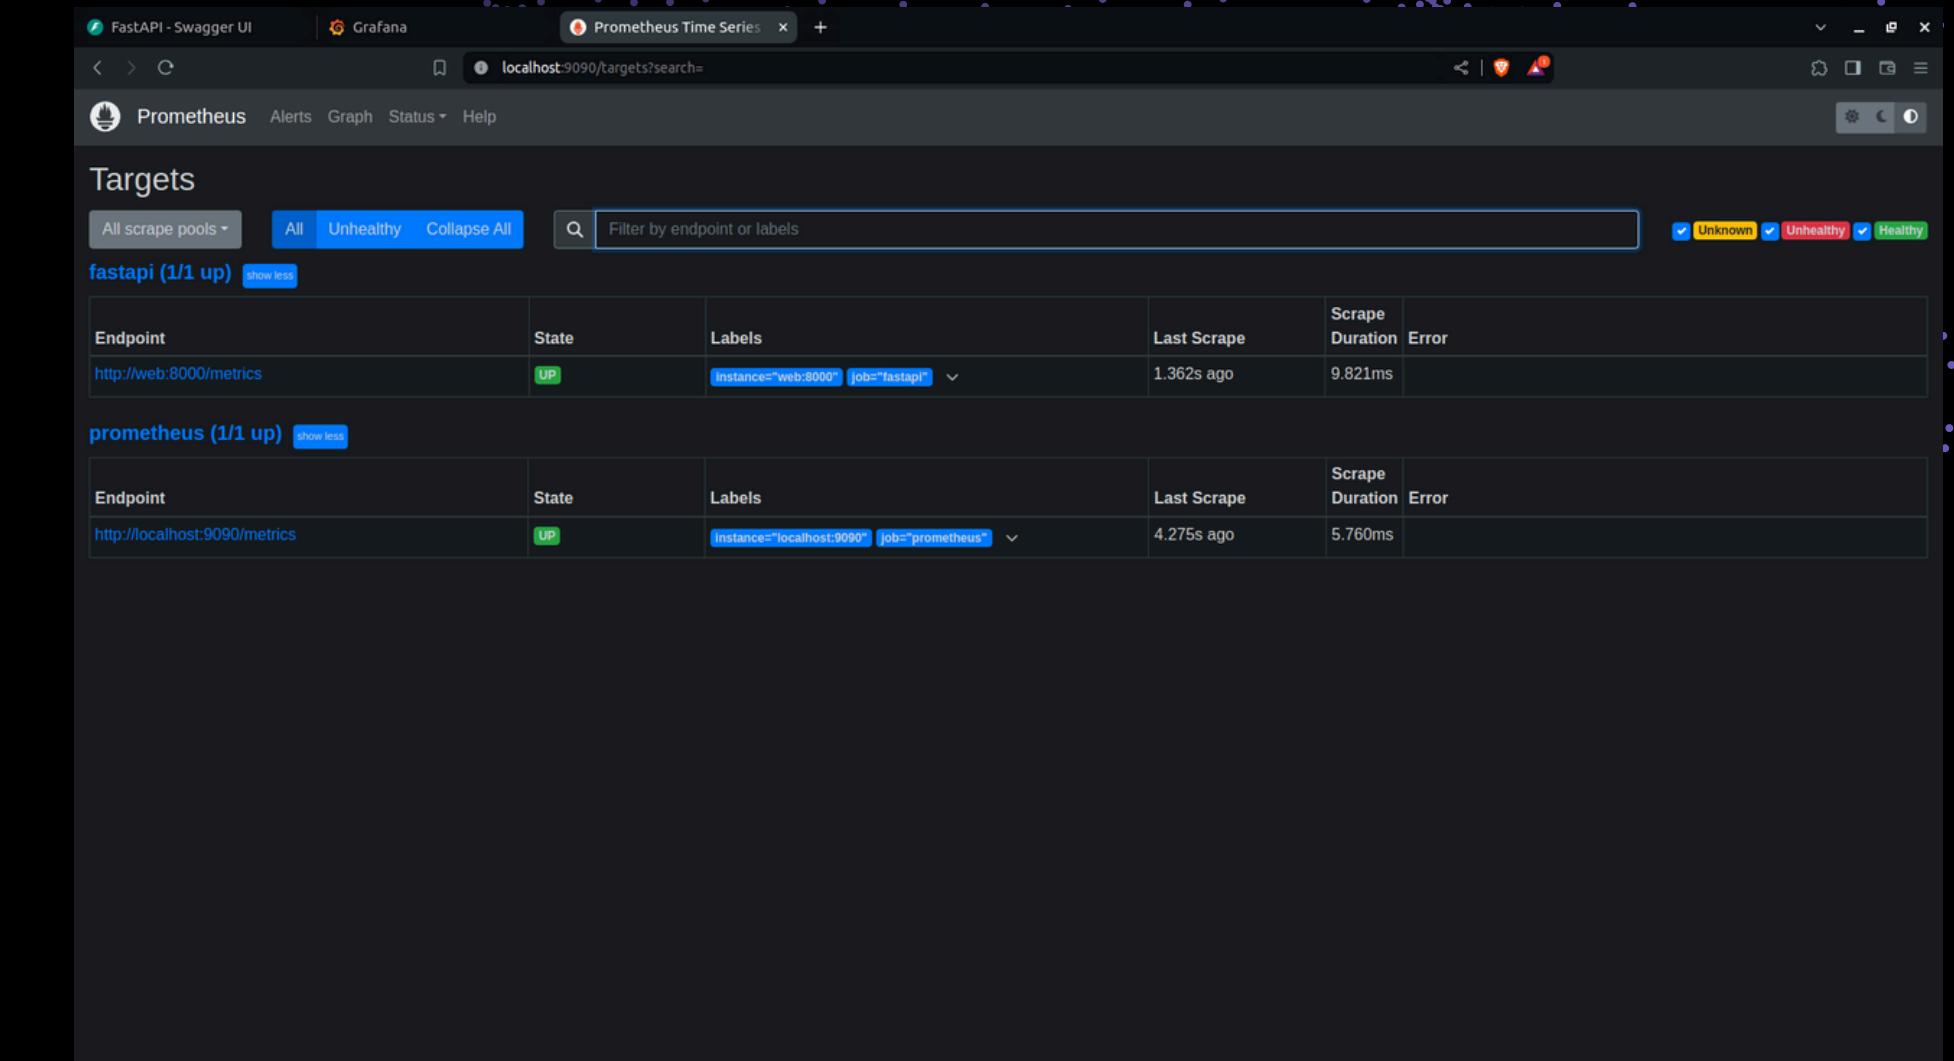
- The best model based on performance on the test data was chosen to be deployed.
- A FastAPI application is designed to host the model's predictive capabilities. The API endpoint accepts image uploads, then processes the image, before passing it through a pretrained CNN model and returns the predicted class label.
- The application is finally hosted using Uvicorn and accessible for SwaggerUI documentation.





Health Monitoring

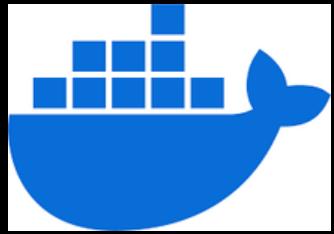
- The application uses the Instrumentator to instrument and expose metrics. Various Prometheus metrics are set up to monitor API usage, run time, memory and CPU usage, as well as network I/O.
- Then customized dashboards have been created using Grafana to visualize and analyze Prometheus metrics for the Rest API. The dashboards empower proactive monitoring and help troubleshooting issues, setting alerts, thus ensuring SLOs are met.





Dashboard





Dockerization

- Docker was used to containerize the FastAPI application, ensuring a consistent and reproducible environment across different stages of development and deployment.
- Alongside FastAPI, Prometheus and Grafana was integrated within the Docker container to collect and store real-time metrics, visualize these metrics, providing a comprehensive dashboard for real-time analytics and insights into the application's performance.
- This containerized setup streamlined the deployment process, enhanced scalability, and ensured that monitoring and visualization tools were readily available and consistently configured across all environments.



Version Control

- Implemented version control using Git with a well structured repository to include data processing scripts, model training code, the FastAPI application, and monitoring setups with Docker configurations.
- Proper documentation detailing methods and usage steps, ensuring each part of the project, from setup to execution has been done using github.

Github repo : <https://github.com/Prithiviraj7R/BDL-Project>

Future Improvements

- Usage of pretrained models like ResNet or GoogLeNet can improve the model performance.
- Train on a larger dataset to have better performance, also with larger dataset, processing would take more time, and hence usage of Spark would become justifiable.

Contributions

R S Harish Krishnan (ME20B081): Integrated Prometheus and Grafana into a FastAPI application and containerized the application using Docker.

R Prithiviraj (ME20B136): Implemented the model and tracked model tuning using MLflow, and developed the FastAPI application.

B Srivathsan (ME20B176): Developed the data processing pipeline using DVC and containerized the application using Docker.

**THANK YOU
SO MUCH!**